

01_eda_pp_vs_oil

December 2, 2025

1 Polypropylene vs Upstream Drivers: Exploratory Data Analysis

Explore whether PP prices co-move with propylene (PGP) and crude; establish sensible baselines before expanding into richer feature sets.

Local files under `data/` are used throughout. Units, currencies, and frequencies may differ, so we normalize and resample to monthly for comparability.

```
[1]: #pip install -q pandas numpy matplotlib seaborn statsmodels
```

```
[2]: from pathlib import Path
import sys
import pandas as pd
import matplotlib.pyplot as plt
from IPython.display import Markdown, display

PROJECT_DIR = Path.cwd().parent
DATA_DIR = PROJECT_DIR / "data"
PLOTS_DIR = PROJECT_DIR / "plots"
ARTIFACTS_DIR = PROJECT_DIR / "artifacts"
for folder in (PLOTS_DIR, ARTIFACTS_DIR):
    folder.mkdir(parents=True, exist_ok=True)

sys.path.insert(0, str(PROJECT_DIR / "src"))
from eda_utils import (
    build_audit_table,
    compute_log_returns,
    cross_corr_scan,
    find_candidate_files,
    load_series,
    normalize_base_100,
    plot_beta_scaled_spread,
    plot_leadlag_heatmap,
    plot_normalized_levels,
    plot_rolling_beta,
    plot_scatter_returns,
    rolling_beta_ols,
    rolling_corr,
```

```

        to_monthly,
    )

plt.rcParams["figure.dpi"] = 150
CLASS_KEYWORDS = {
    "PP": ["pp", "polypropylene"],
    "PGP": ["pgp", "propylene", "c3"],
    "CRUDE": ["brent", "wti", "crude", "oil"],
    "NAPHTHA": ["naphtha"],
}

```

1.1 Discovery & loading

Find candidate files by keyword, load each series with `eda_utils.load_series`, and capture metadata for traceability.

```
[3]: candidate_map = {
    label: find_candidate_files(DATA_DIR, patterns)
    for label, patterns in CLASS_KEYWORDS.items()
}

loaded_series = []
discovery_records = []
for label, paths in candidate_map.items():
    if not paths:
        print(
            f"No {label} series found. Expected keywords: "
            f"{CLASS_KEYWORDS[label]}"
        )
        continue
    series_list = []
    for path in paths:
        try:
            df = load_series(path, f"{label}|{path.stem}")
        except Exception as exc:
            print(f"Skipping {path.name}: {exc}")
            continue
        series_list.append(df)
        discovery_records.append(
            {
                "class": label,
                "series_name": df.attrs.get("series_name"),
                "file": path.relative_to(PROJECT_DIR),
                "date_column": df.attrs.get("date_column"),
                "value_column": df.attrs.get("value_column"),
                "rows_raw": df.attrs.get("rows_raw"),
                "rows_clean": df.attrs.get("rows_clean"),
            }
        )
}
```

```

        "start": df.attrs.get("start"),
        "end": df.attrs.get("end"),
    }
)
if series_list:
    loaded_series[label] = series_list
discovery_df = pd.DataFrame(discovery_records)
if not discovery_df.empty:
    display(discovery_df)
else:
    display(Markdown("No qualifying series loaded - add data to proceed."))

```

No NAPHTHA series found. Expected keywords: ['naphtha']

	class	series_name \
0	PP	PP polypropylene_primary_avg_prices
1	PP	PP polypropylene_weekly
2	PP	PP polypropylene_weekly_clean
3	PGP	PGP polypropylene_primary_avg_prices
4	PGP	PGP polypropylene_weekly
5	PGP	PGP polypropylene_weekly_clean
6	CRUDE	CRUDE crude_oil_daily
7	CRUDE	CRUDE crude_oil_weekly
8	CRUDE	CRUDE crude_oil_weekly_clean

	file	date_column \
0	data\prices\polypropylene_primary_avg_prices.csv	Date
1	data\prices\polypropylene_weekly.csv	Date
2	data\prices\polypropylene_weekly_clean.csv	date
3	data\prices\polypropylene_primary_avg_prices.csv	Date
4	data\prices\polypropylene_weekly.csv	Date
5	data\prices\polypropylene_weekly_clean.csv	date
6	data\prices\crude_oil_daily.csv	Date
7	data\prices\crude_oil_weekly.csv	Date
8	data\prices\crude_oil_weekly_clean.csv	Date

	value_column	rows_raw	rows_clean	start	end
0	PP_Avg_EUR_per_t	81	81	2018-08-01	2025-04-01
1	Price	520	520	2015-01-04	2025-03-02
2	price	520	520	2015-01-04	2025-03-02
3	PP_Avg_EUR_per_t	81	81	2018-08-01	2025-04-01
4	Price	520	520	2015-01-04	2025-03-02
5	price	520	520	2015-01-04	2025-03-02
6	Close	5426	2712	2015-01-02	2025-10-14
7	Close	1129	564	2015-01-05	2025-10-20
8	Close	1128	1128	2015-01-05	2025-10-20

Table note: The discovery summary above lists every file that matched our keyword scan and highlights which date/value columns were inferred so you can quickly spot mislabelled sources.

1.2 Monthly alignment

Resample each raw series to monthly means, keep one preferred series per class, and form wide outer/inner tables.

```
[4]: monthly_tables = {}
preferred = {}
for label, series_list in loaded_series.items():
    monthly_list = [to_monthly(series) for series in series_list]
    table = pd.concat(monthly_list, axis=1)
    table = table.loc[~table.index.duplicated()].sort_index()
    monthly_tables[label] = table
    preferred_series_name = table.count().sort_values(ascending=False).index[0]
    preferred[label] = table[preferred_series_name]
wide_outer = pd.concat(monthly_tables.values(), axis=1) if monthly_tables else pd.DataFrame()
monthly_inner = pd.concat(preferred, axis=1).dropna() if preferred else pd.DataFrame()
display(Markdown(f"**Monthly outer shape:** {wide_outer.shape}"))
display(Markdown(f"**Monthly inner shape:** {monthly_inner.shape}"))
wide_outer.head()
```

Monthly outer shape: (130, 9)

Monthly inner shape: (123, 3)

```
[4]:          PP|polypropylene_primary_avg_prices  PP|polypropylene_weekly \
2015-01-01                      NaN            8285.25
2015-02-01                      NaN            8748.00
2015-03-01                      NaN            8829.40
2015-04-01                      NaN            9348.00
2015-05-01                      NaN            9142.00

          PP|polypropylene_weekly_clean \
2015-01-01            8285.25
2015-02-01            8748.00
2015-03-01            8829.40
2015-04-01            9348.00
2015-05-01            9142.00

          PGP|polypropylene_primary_avg_prices  PGP|polypropylene_weekly \
2015-01-01                      NaN            8285.25
2015-02-01                      NaN            8748.00
2015-03-01                      NaN            8829.40
2015-04-01                      NaN            9348.00
2015-05-01                      NaN            9142.00

          PGP|polypropylene_weekly_clean  CRUDE|crude_oil_daily \
2015-01-01            8285.25            47.325500
```

2015-02-01	8748.00	50.724736
2015-03-01	8829.40	47.854091
2015-04-01	9348.00	54.628096
2015-05-01	9142.00	59.372000

	CRUDE crude_oil_weekly	CRUDE crude_oil_weekly_clean
2015-01-01	48.224125	49.246187
2015-02-01	49.986625	53.093875
2015-03-01	48.135600	52.875400
2015-04-01	53.463125	56.592562
2015-05-01	59.418000	62.724000

```
[5]: audit_input = {label: series for label, series in preferred.items()}
audit_table = build_audit_table(audit_input) if audit_input else pd.DataFrame()
if not audit_table.empty:
    display(audit_table)
    audit_table.to_csv(ARTIFACTS_DIR / "data_audit.csv", index=False)
monthly_inner.to_csv(ARTIFACTS_DIR / "merged_monthly_prices.csv")
```

	series	first_date	last_date	count	missing_pct
0	PP	2015-01-01	2025-03-01	123	0.0
1	PGP	2015-01-01	2025-03-01	123	0.0
2	CRUDE	2015-01-01	2025-10-01	130	0.0

	frequency_note	currency	unit	source_path
0	monthly (resampled via mean)	None	None	None
1	monthly (resampled via mean)	None	None	None
2	monthly (resampled via mean)	None	None	None

Audit note: Use the audit table to confirm coverage (first/last date, missing %), plus currency/unit hints pulled from the raw files.

1.3 Levels (normalized)

Normalize each preferred level series to 100 at the first common observation.

```
[6]: if not monthly_inner.empty:
    norm_levels = normalize_base_100(monthly_inner)
    fig = plot_normalized_levels(norm_levels, "Normalized Levels (base = 100)")
    display(fig)
    fig.savefig(PLOTS_DIR / "levels_normalized.png", bbox_inches="tight")
    plt.close(fig)
else:
    display(Markdown("*No overlapping monthly series available.*"))
```




Figure: Normalized level lines let us compare PP and crude trends on a common 100-base scale so different currencies/units do not mask co-movement.

1.4 Returns & correlations

Compute log returns, save clean datasets, and prepare diagnostics for PP vs key drivers.

```
[7]: if not monthly_inner.empty:
    log_returns = compute_log_returns(monthly_inner).dropna()
    log_returns.to_csv(ARTIFACTS_DIR / "merged_monthly_returns.csv")
    display(log_returns.describe().T)
else:
    log_returns = pd.DataFrame()
    display(Markdown("*Returns unavailable; add overlapping series.*"))
```

	count	mean	std	min	25%	50%	75%	\
PP	122.0	-0.001052	0.042444	-0.131583	-0.023175	-0.001969	0.020649	
PGP	122.0	-0.001052	0.042444	-0.131583	-0.023175	-0.001969	0.020649	
CRUDE	122.0	0.002964	0.120098	-0.600585	-0.053332	0.019625	0.056092	

	max
PP	0.149596
PGP	0.149596
CRUDE	0.535517

Table detail: The return summary offers mean/volatility/skewness for each monthly log-return series so you can judge distribution shape before modelling.

1.5 Headline visuals

Key diagnostics for PP versus crude oil (or other drivers when available).

```
[8]: if {"PP", "CRUDE"} <= set(monthly_inner.columns):
    pp = monthly_inner["PP"]
    crude = monthly_inner["CRUDE"]
    pp_ret = log_returns["PP"]
    crude_ret = log_returns["CRUDE"]



    rolling_beta = rolling_beta_ols(pp_ret, crude_ret, window=12)
    if not rolling_beta.empty:
        fig = plot_rolling_beta(rolling_beta, "Rolling 12M OLS : PP vs Crude")
        display(fig)
        fig.savefig(PLOTS_DIR / "rolling_beta_pp_crude.png", bbox_inches="tight")
        plt.close(fig)



    beta_series = rolling_beta["beta"] if "beta" in rolling_beta else None
    fig = plot_beta_scaled_spread(pp, crude, beta=None, rolling_beta=beta_series, title="Beta-scaled spread: PP vs Crude")
    display(fig)
    fig.savefig(PLOTS_DIR / "beta_scaled_spread_pp_crude.png", bbox_inches="tight")
    plt.close(fig)

    fig = plot_leadlag_heatmap(pp_ret, crude_ret, window=36, max_lag=12, title="Lead/Lag heatmap: PP vs Crude")
    display(fig)
    fig.savefig(PLOTS_DIR / "leadlag_heatmap_pp_crude.png", bbox_inches="tight")
    plt.close(fig)

    fig = plot_scatter_returns(pp_ret, crude_ret, title="Monthly log returns: PP vs Crude", label_outliers=3)
    display(fig)
    fig.savefig(PLOTS_DIR / "scatter_returns_pp_crude.png", bbox_inches="tight")
    plt.close(fig)

    rolling_corr_pp_crude = rolling_corr(pp_ret, crude_ret, window=12)
    rolling_corr_pp_crude.to_csv(ARTIFACTS_DIR / "rolling_corr_pp_crude.csv")
    corr_scan = cross_corr_scan(pp_ret, crude_ret, max_lag=12)
    corr_scan.to_csv(ARTIFACTS_DIR / "cross_corr_scan_pp_crude.csv", index=False)
else:
    display(Markdown("Add crude data to unlock headline visuals.**"))
```





Figures recap: Rolling tracks relationship strength, the beta-scaled spread highlights sustained divergence, the lead/lag heatmap surfaces timing asymmetry, and the scatter pinpoints outlier months that may warrant narrative follow-up.

1.6 Findings

```
[9]: if not monthly_inner.empty and not log_returns.empty:
    overlap = (
        monthly_inner.index.min().date(),
        monthly_inner.index.max().date(),
```

```

    len(monthly_inner),
)
corr_pp_crude = None
if {"PP", "CRUDE"} <= set(log_returns.columns):
    corr_pp_crude = log_returns[["PP", "CRUDE"]].corr().iloc[0, 1]
bullets = [
    f"- Overlap: {overlap[0]} to {overlap[1]} ({overlap[2]} months).",
    f"- PP vs Crude return correlation: {corr_pp_crude:.2f}." if_
    ↪corr_pp_crude is not None else "- Crude returns unavailable.",
]
if Path(ARTIFACTS_DIR / "cross_corr_scan_pp_crude.csv").exists():
    scan = pd.read_csv(ARTIFACTS_DIR / "cross_corr_scan_pp_crude.csv")
    best = scan.iloc[scan["corr"].abs().idxmax()]
    bullets.append(
        f"- Best lead/lag (PP vs Crude): lag {int(best['lag'])} with corr_"
    ↪{best['corr']:.2f}."
    )
else:
    bullets.append("- Add drivers (PGP/crude) to compute lead/lag"
    ↪diagnostics.)
display(Markdown("\n".join(bullets)))
else:
    display(Markdown("**Not enough data for findings yet.**"))

```

- Overlap: 2015-01-01 to 2025-03-01 (123 months).
- PP vs Crude return correlation: 0.21.
- Best lead/lag (PP vs Crude): lag -1 with corr 0.36.