

07_sentiment_modeling_dataset

December 2, 2025

1 Sentiment Modeling Dataset

Build a single monthly dataset that combines PP prices/returns, PP residuals, and OPEC sentiment indices, then run a quick EDA to understand sentiment vs PP before modeling. All inputs are precomputed artifacts; no external APIs are used.

```
[1]: from pathlib import Path
import sys
import warnings

import pandas as pd
import matplotlib.pyplot as plt
from IPython.display import display

NOTEBOOK_DIR = Path(__file__).resolve().parent if "__file__" in globals() else Path.cwd()
ROOT = NOTEBOOK_DIR.parent
if str(ROOT) not in sys.path:
    sys.path.append(str(ROOT))

from src.model_dataset_utils import (
    build_modeling_dataset,
    summarize_modeling_dataset,
    plot_pp_vs_opec_index,
    plot_residual_vs_index_scatter,
    plot_leadlag_sentiment_vs_residual,
)

DATA = ROOT / "data"
ART = ROOT / "artifacts"
PLOTS = ROOT / "plots"
for path in (ART, PLOTS):
    path.mkdir(parents=True, exist_ok=True)

warnings.filterwarnings("ignore")
pd.set_option("display.max_columns", 40)
print(f"ROOT set to {ROOT}")
```

```
ROOT set to c:\PythonProjects\LLM-polypropylene
```

```
[2]: df_model = pd.DataFrame()
meta = {}

try:
    df_model, meta = build_modeling_dataset(ROOT)
except Exception as exc:
    print(f"Failed to build modeling dataset: {exc}")

if df_model.empty:
    print("Modeling dataset is empty; downstream steps will be skipped.")
else:
    print(f"Shape: {df_model.shape}")
    print(f"Date range: {df_model.index.min().date()} to {df_model.index.max().date()}")
    print(f"Columns: {list(df_model.columns)}")
    display(df_model.head())
    display(df_model.tail())
    print("Metadata:", meta)
```

```
Loaded monthly prices from c:\PythonProjects\LLM-
polypropylene\artifacts\merged_monthly_prices.csv, columns: ['PP', 'PGP',
'CRUDE'], coverage: 2015-01-31 to 2025-03-31
Loaded monthly returns from c:\PythonProjects\LLM-
polypropylene\artifacts\merged_monthly_returns.csv, columns: ['PP', 'PGP',
'CRUDE'], coverage: 2015-02-28 to 2025-03-31
Loaded residuals from c:\PythonProjects\LLM-
polypropylene\artifacts\pp_idiosyncratic_residual.csv, rows: 119, coverage:
2015-05-31 to 2025-03-31
Loaded OPEC monthly index from c:\PythonProjects\LLM-
polypropylene\artifacts\opec_sentiment_monthly_index.csv, columns:
['score_demand', 'score_overall', 'score_price_outlook', 'score_supply',
'index_demand', 'index_overall', 'index_price_outlook', 'index_supply',
'index_hybrid']
Loaded OPEC section scores from c:\PythonProjects\LLM-
polypropylene\artifacts\opec_sentiment_monthly_section_scores.csv
OPEC index coverage: 2019-01-31 to 2025-11-30
Modeling dataset shape: (130, 16)
Coverage: 2015-02-28 to 2025-11-30
Shape: (130, 16)
Date range: 2015-02-28 to 2025-11-30
Columns: ['PP', 'PGP', 'CRUDE', 'ret_PP', 'ret_PGP', 'ret_CRUDE', 'resid_PP',
'score_demand', 'score_overall', 'score_price_outlook', 'score_supply',
'index_demand', 'index_overall', 'index_price_outlook', 'index_supply',
'index_hybrid']
```

```
PP      PGP      CRUDE      ret_PP   ret_PGP  ret_CRUDE  \
```

2015-02-28	8748.0	8748.0	50.724736	0.054348	0.054348	0.069364		
2015-03-31	8829.4	8829.4	47.854091	0.009262	0.009262	-0.058257		
2015-04-30	9348.0	9348.0	54.628096	0.057075	0.057075	0.132392		
2015-05-31	9142.0	9142.0	59.372000	-0.022283	-0.022283	0.083274		
2015-06-30	8786.0	8786.0	59.828637	-0.039720	-0.039720	0.007662		
	resid_PP	score_demand	score_overall	score_price_outlook			\	
2015-02-28	NaN		NaN		NaN		NaN	
2015-03-31	NaN		NaN		NaN		NaN	
2015-04-30	NaN		NaN		NaN		NaN	
2015-05-31	-0.025787		NaN		NaN		NaN	
2015-06-30	-0.033245		NaN		NaN		NaN	
	score_supply	index_demand	index_overall	index_price_outlook			\	
2015-02-28	NaN		NaN		NaN		NaN	
2015-03-31	NaN		NaN		NaN		NaN	
2015-04-30	NaN		NaN		NaN		NaN	
2015-05-31	NaN		NaN		NaN		NaN	
2015-06-30	NaN		NaN		NaN		NaN	
	index_supply	index_hybrid						
2015-02-28	NaN		NaN					
2015-03-31	NaN		NaN					
2015-04-30	NaN		NaN					
2015-05-31	NaN		NaN					
2015-06-30	NaN		NaN					
	PP	PGP	CRUDE	ret_PP	ret_PGP	ret_CRUDE	resid_PP	\
2025-07-31	NaN	NaN	NaN	NaN	NaN	NaN	NaN	
2025-08-31	NaN	NaN	NaN	NaN	NaN	NaN	NaN	
2025-09-30	NaN	NaN	NaN	NaN	NaN	NaN	NaN	
2025-10-31	NaN	NaN	NaN	NaN	NaN	NaN	NaN	
2025-11-30	NaN	NaN	NaN	NaN	NaN	NaN	NaN	
	score_demand	score_overall	score_price_outlook	score_supply			\	
2025-07-31	-0.465910		0.7	0.730664		-0.382759		
2025-08-31	-0.439590		0.7	0.341918		-0.376988		
2025-09-30	-0.592581		0.3	0.675067		-0.389453		
2025-10-31	-0.494205		0.4	0.173203		-0.468063		
2025-11-30	-0.406237		0.3	-0.027499		-0.509039		
	index_demand	index_overall	index_price_outlook	index_supply			\	
2025-07-31	-13.950573		8.3	33.377305		5.704018		
2025-08-31	-14.390164		9.0	33.719222		5.327030		
2025-09-30	-14.982745		9.3	34.394289		4.937577		
2025-10-31	-15.476949		9.7	34.567492		4.469514		
2025-11-30	-15.883186		10.0	34.539993		3.960475		

```

    index_hybrid
2025-07-31      8.357687
2025-08-31      8.414022
2025-09-30      8.412280
2025-10-31      8.315014
2025-11-30      8.154320

Metadata: {'start_date': '2015-02-28', 'end_date': '2025-11-30', 'n_months': '130', 'columns': 'PP, PGP, CRUDE, ret_PP, ret_PGP, ret_CRUDE, resid_PP, score_demand, score_overall, score_price_outlook, score_supply, index_demand, index_overall, index_price_outlook, index_supply, index_hybrid'}

```

```
[3]: if df_model.empty:
    print("Skipping dataset export; nothing to save.")
else:
    out_path = ART / "modeling_dataset_monthly.csv"
    df_model.reset_index().rename(columns={"index": "date"}).to_csv(out_path, index=False)
    print(f"Saved modeling dataset -> {out_path}")

```

Saved modeling dataset -> c:\PythonProjects\LLM-polypropylene\artifacts\modeling_dataset_monthly.csv

```
[4]: summaries = summarize_modeling_dataset(df_model) if not df_model.empty else {}

if not summaries or df_model.empty:
    print('No summaries to display.')
else:
    desc = summaries.get('describe')
    miss = summaries.get('missing')
    corr = summaries.get('corr')

    print('Descriptive stats:')
    display(desc)

    print("\nMissing values:")
    display(miss)

    if corr is not None and not corr.empty:
        print("\nCorrelation matrix:")
        display(corr)

# Save summaries
desc_path = ART / 'modeling_dataset_describe.csv'
miss_path = ART / 'modeling_dataset_missing.csv'
corr_path = ART / 'modeling_dataset_corr.csv'
desc.to_csv(desc_path)
miss.to_csv(miss_path)
```

```

if corr is not None and not corr.empty:
    corr.to_csv(corr_path)
    print(f'Saved summaries to {ART}')

```

Descriptive stats:

	count	mean	std	min	\
PP	122.0	8.190967e+03	855.627555	6043.600000	
PGP	122.0	8.190967e+03	855.627555	6043.600000	
CRUDE	122.0	6.232619e+01	17.816844	16.699048	
ret_PP	122.0	-1.052335e-03	0.042444	-0.131583	
ret_PGP	122.0	-1.052335e-03	0.042444	-0.131583	
ret_CRUDE	122.0	2.963678e-03	0.120098	-0.600585	
resid_PP	119.0	1.012262e-16	0.041258	-0.121069	
score_demand	82.0	-1.896560e-01	0.306369	-0.592581	
score_overall	81.0	1.234568e-01	0.560641	-0.900000	
score_price_outlook	82.0	4.249247e-01	0.314560	-0.098708	
score_supply	82.0	4.798564e-02	0.345477	-0.509039	
index_demand	83.0	-4.035231e+00	5.850803	-15.883186	
index_overall	83.0	-3.144578e-01	5.959500	-10.800000	
index_price_outlook	83.0	1.613468e+01	10.261895	0.000000	
index_supply	83.0	6.833947e+00	3.717035	0.000000	
index_hybrid	83.0	4.654735e+00	3.228771	-0.310274	

	25%	50%	75%	max
PP	7566.737500	8186.000000	8800.062500	10339.250000
PGP	7566.737500	8186.000000	8800.062500	10339.250000
CRUDE	49.058365	60.347045	73.738866	114.339048
ret_PP	-0.023175	-0.001969	0.020649	0.149596
ret_PGP	-0.023175	-0.001969	0.020649	0.149596
ret_CRUDE	-0.053332	0.019625	0.056092	0.535517
resid_PP	-0.023786	-0.001608	0.017806	0.135678
score_demand	-0.436936	-0.216037	0.022726	0.532321
score_overall	-0.400000	0.400000	0.600000	0.800000
score_price_outlook	0.166677	0.322499	0.748457	0.980334
score_supply	-0.345003	0.080297	0.162933	0.658455
index_demand	-9.191848	-2.458879	0.645331	3.228908
index_overall	-5.350000	-0.700000	5.300000	10.000000
index_price_outlook	6.975030	14.961675	24.227515	34.567492
index_supply	3.930363	7.593114	10.350194	12.371179
index_hybrid	0.881179	5.296614	7.923849	8.593474

Missing values:

	missing_count
PP	8
PGP	8
CRUDE	8

ret_PP	8
ret_PGP	8
ret_CRUDE	8
resid_PP	11
score_demand	48
score_overall	49
score_price_outlook	48
score_supply	48
index_demand	47
index_overall	47
index_price_outlook	47
index_supply	47
index_hybrid	47

Correlation matrix:

	ret_PP	resid_PP	index_hybrid	index_demand	\
ret_PP	1.000000	9.749861e-01	0.058252	0.031054	
resid_PP	0.974986	1.000000e+00	0.066762	0.007338	
index_hybrid	0.058252	6.676164e-02	1.000000	-0.868254	
index_demand	0.031054	7.338081e-03	-0.868254	1.000000	
index_supply	0.097761	9.714554e-02	0.811000	-0.459749	
index_price_outlook	0.073303	8.668097e-02	0.953841	-0.941189	
ret_CRUDE	0.209743	-1.958379e-16	-0.023961	0.099182	
ret_PGP	1.000000	9.749861e-01	0.058252	0.031054	
	index_supply	index_price_outlook	ret_CRUDE	ret_PGP	
ret_PP	0.097761	0.073303	2.097435e-01	1.000000	
resid_PP	0.097146	0.086681	-1.958379e-16	0.974986	
index_hybrid	0.811000	0.953841	-2.396133e-02	0.058252	
index_demand	-0.459749	-0.941189	9.918233e-02	0.031054	
index_supply	1.000000	0.627322	1.533391e-02	0.097761	
index_price_outlook	0.627322	1.000000	-4.561815e-02	0.073303	
ret_CRUDE	0.015334	-0.045618	1.000000e+00	0.209743	
ret_PGP	0.097761	0.073303	2.097435e-01	1.000000	

Saved summaries to c:\PythonProjects\LLM-polypropylene\artifacts

```
[5]: if df_model.empty:
    print("Skipping plots; dataset is empty.")
else:
    # Time-series: PP residuals/returns vs hybrid index
    try:
        fig_ts = plot_pp_vs_opec_index(df_model)
        path_ts = PLOTS / "model_pp_vs_opec_index_timeseries.png"
        fig_ts.savefig(path_ts, dpi=150, bbox_inches="tight")
        display(fig_ts)
    except Exception as exc:
```

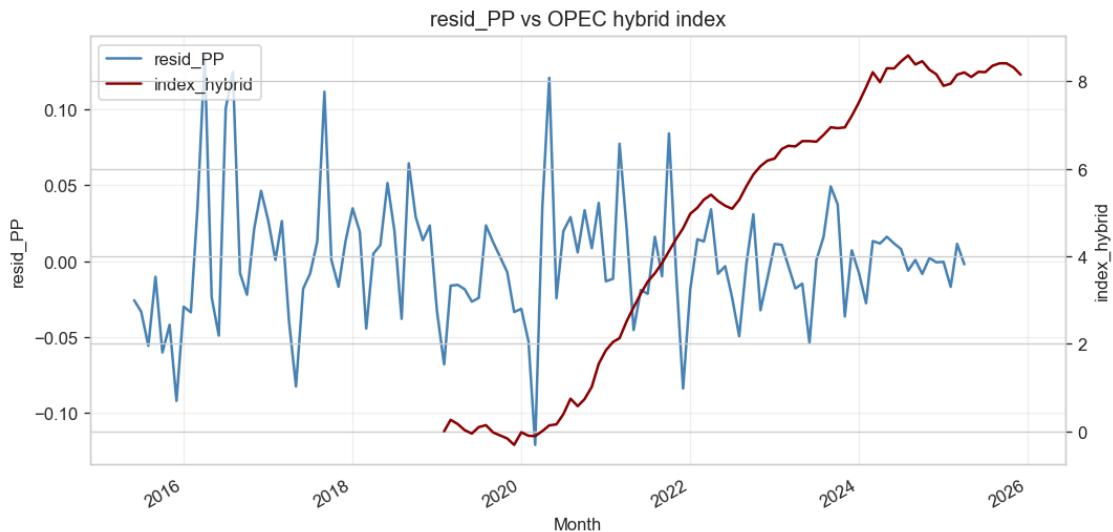
```

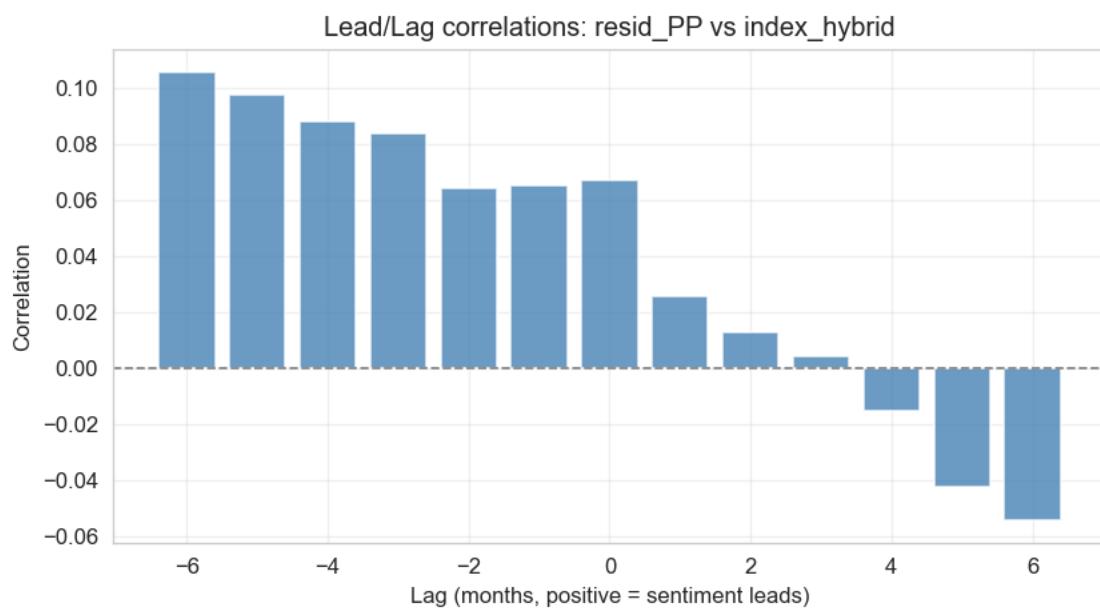
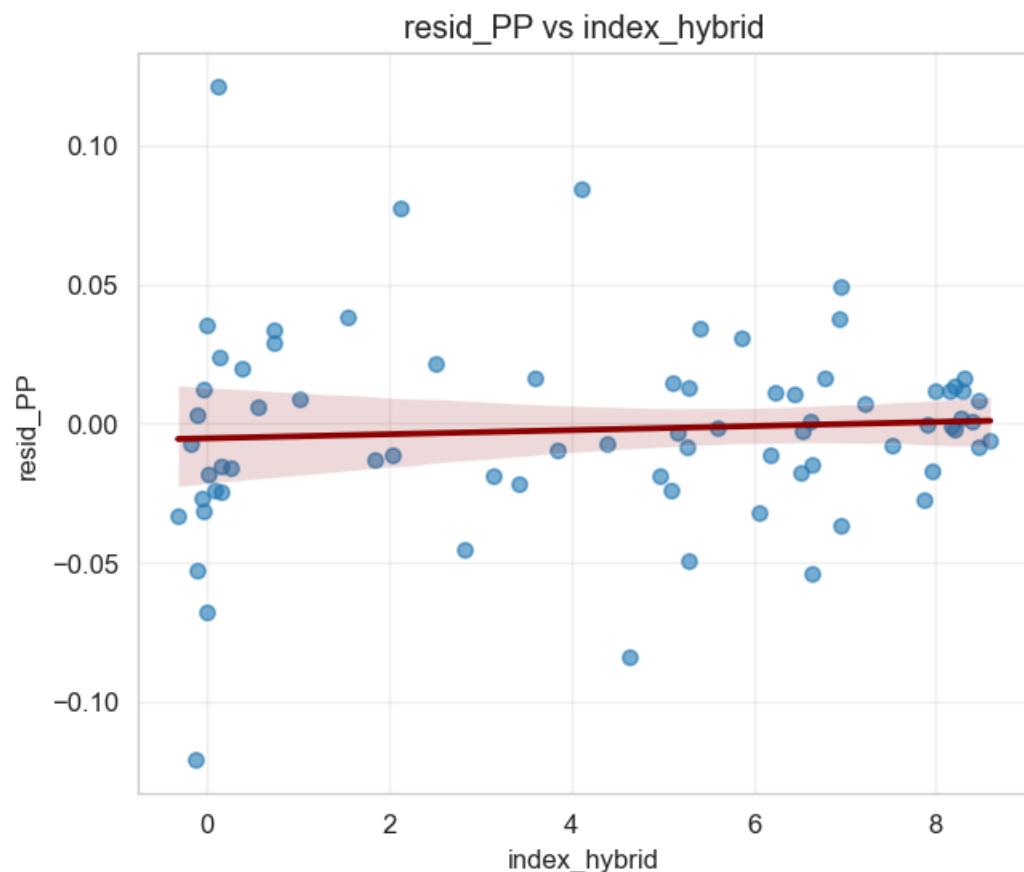
        print(f"Time-series plot failed: {exc}")
    finally:
        plt.close('all')

# Scatter residuals vs index
try:
    fig_scatter = plot_residual_vs_index_scatter(df_model)
    path_scatter = PLOTS / "model_pp_residual_vs_opec_index_scatter.png"
    fig_scatter.savefig(path_scatter, dpi=150, bbox_inches="tight")
    display(fig_scatter)
except Exception as exc:
    print(f"Scatter plot failed: {exc}")
finally:
    plt.close('all')

# Lead/lag correlations
try:
    fig_leadlag = plot_leadlag_sentiment_vs_residual(df_model, max_lag=6)
    path_leadlag = PLOTS / "model_leadlag_residual_index_hybrid.png"
    fig_leadlag.savefig(path_leadlag, dpi=150, bbox_inches="tight")
    display(fig_leadlag)
except Exception as exc:
    print(f"Lead/lag plot failed: {exc}")
finally:
    plt.close('all')

```





```
[6]: if df_model.empty:
    print("No dataset to interpret.")
else:
    start = df_model.index.min().date()
    end = df_model.index.max().date()
    print(f"Modeling dataset spans {start} to {end} with {len(df_model)} months.
         ")
    print("Use these merged features (PP, returns, residuals, OPEC indices) in\u202a
         the next notebook for sentiment-augmented forecasting.")
```

Modeling dataset spans 2015-02-28 to 2025-11-30 with 130 months.

Use these merged features (PP, returns, residuals, OPEC indices) in the next notebook for sentiment-augmented forecasting.