

# 08\_sentiment\_forecasting\_models

December 2, 2025

## 1 Sentiment-Augmented Forecasting

Use the combined modeling dataset to compare baseline PP models vs sentiment-augmented models. We run rolling-origin tests, compute metrics, and inspect coefficients to quantify the incremental value of OPEC sentiment.

```
[1]: from pathlib import Path
import sys
import warnings

import pandas as pd
import matplotlib.pyplot as plt
from IPython.display import display

NOTEBOOK_DIR = Path(__file__).resolve().parent if "__file__" in globals() else Path.cwd()
ROOT = NOTEBOOK_DIR.parent
if str(ROOT) not in sys.path:
    sys.path.append(str(ROOT))

from src.model_dataset_utils import build_modeling_dataset
from src.sentiment_model_utils import (
    define_model_specs,
    run_models_backtest,
    compute_forecast_metrics,
    extract_coefficients,
    plot_forecast_vs_actual,
    plot_forecast_errors,
    plot_coefficients,
)
DATA = ROOT / "data"
ART = ROOT / "artifacts"
PLOTS = ROOT / "plots"
for path in (ART, PLOTS):
    path.mkdir(parents=True, exist_ok=True)

warnings.filterwarnings("ignore")
```

```
pd.set_option("display.max_columns", 50)
print(f"ROOT set to {ROOT}")
```

ROOT set to c:\PythonProjects\LLM-polypropylene

```
[2]: df_model = pd.DataFrame()

artifact_path = ART / "modeling_dataset_monthly.csv"
if artifact_path.exists():
    df_model = pd.read_csv(artifact_path, parse_dates=["date"])
    df_model = df_model.set_index("date").sort_index()
    print(f"Loaded modeling dataset from {artifact_path}")
else:
    print("Modeling dataset artifact missing; rebuilding in memory via"
          "build_modeling_dataset().")
    df_model, _ = build_modeling_dataset(ROOT)

if df_model.empty:
    print("Modeling dataset is empty; please run notebooks 02, 06, 07 first.")
else:
    print(f"Shape: {df_model.shape}")
    print(f"Date range: {df_model.index.min().date()} to {df_model.index.max().date()}")
    print(f"Columns: {[list(df_model.columns)}")
    display(df_model.head())
    display(df_model.tail())
```

Loaded modeling dataset from c:\PythonProjects\LLM-polypropylene\artifacts\modeling\_dataset\_monthly.csv

Shape: (130, 16)

Date range: 2015-02-28 to 2025-11-30

Columns: ['PP', 'PGP', 'CRUDE', 'ret\_PP', 'ret\_PGP', 'ret\_CRUDE', 'resid\_PP', 'score\_demand', 'score\_overall', 'score\_price\_outlook', 'score\_supply', 'index\_demand', 'index\_overall', 'index\_price\_outlook', 'index\_supply', 'index\_hybrid']

	PP	PGP	CRUDE	ret_PP	ret_PGP	ret_CRUDE	\
--	----	-----	-------	--------	---------	-----------	---

date

2015-02-28	8748.0	8748.0	50.724736	0.054348	0.054348	0.069364	
2015-03-31	8829.4	8829.4	47.854091	0.009262	0.009262	-0.058257	
2015-04-30	9348.0	9348.0	54.628096	0.057075	0.057075	0.132392	
2015-05-31	9142.0	9142.0	59.372000	-0.022283	-0.022283	0.083274	
2015-06-30	8786.0	8786.0	59.828637	-0.039720	-0.039720	0.007662	

	resid_PP	score_demand	score_overall	score_price_outlook	\
--	----------	--------------	---------------	---------------------	---

date

2015-02-28	NaN	NaN	NaN	NaN
2015-03-31	NaN	NaN	NaN	NaN
2015-04-30	NaN	NaN	NaN	NaN

2015-05-31	-0.025787	NaN	NaN	NaN				
2015-06-30	-0.033245	NaN	NaN	NaN				
		score_supply	index_demand	index_overall	index_price_outlook	\		
date								
2015-02-28		NaN	NaN	NaN	NaN	NaN		
2015-03-31		NaN	NaN	NaN	NaN	NaN		
2015-04-30		NaN	NaN	NaN	NaN	NaN		
2015-05-31		NaN	NaN	NaN	NaN	NaN		
2015-06-30		NaN	NaN	NaN	NaN	NaN		
		index_supply	index_hybrid					
date								
2015-02-28		NaN	NaN					
2015-03-31		NaN	NaN					
2015-04-30		NaN	NaN					
2015-05-31		NaN	NaN					
2015-06-30		NaN	NaN					
	PP	PGP	CRUDE	ret_PP	ret_PGP	ret_CRUDE	resid_PP	\
date								
2025-07-31	NaN	NaN	NaN	NaN	NaN	NaN	NaN	
2025-08-31	NaN	NaN	NaN	NaN	NaN	NaN	NaN	
2025-09-30	NaN	NaN	NaN	NaN	NaN	NaN	NaN	
2025-10-31	NaN	NaN	NaN	NaN	NaN	NaN	NaN	
2025-11-30	NaN	NaN	NaN	NaN	NaN	NaN	NaN	
		score_demand	score_overall	score_price_outlook	score_supply	\		
date								
2025-07-31	-0.465910		0.7		0.730664		-0.382759	
2025-08-31	-0.439590		0.7		0.341918		-0.376988	
2025-09-30	-0.592581		0.3		0.675067		-0.389453	
2025-10-31	-0.494205		0.4		0.173203		-0.468063	
2025-11-30	-0.406237		0.3		-0.027499		-0.509039	
		index_demand	index_overall	index_price_outlook	index_supply	\		
date								
2025-07-31	-13.950573		8.3		33.377305		5.704018	
2025-08-31	-14.390164		9.0		33.719222		5.327030	
2025-09-30	-14.982745		9.3		34.394289		4.937577	
2025-10-31	-15.476949		9.7		34.567492		4.469514	
2025-11-30	-15.883186		10.0		34.539993		3.960475	
		index_hybrid						
date								
2025-07-31	8.357687							
2025-08-31	8.414022							
2025-09-30	8.412280							

```
2025-10-31      8.315014
2025-11-30      8.154320
```

```
[3]: if df_model.empty:
    print("No data to define model specs.")
else:
    specs = define_model_specs(df_model)
    if not specs:
        print("No model specs could be defined (missing necessary columns).")
    else:
        for s in specs:
            print(f"Model {s.name}: target={s.target}, regressors={s.
→regressors}, target_lags={s.target_lags}, reg_lags={s.reg_lags}")
```

```
Model A_baseline: target=resid_PP, regressors=['ret_CRUDE', 'ret_PGP'],
target_lags=[1], reg_lags={'ret_CRUDE': [1], 'ret_PGP': [1]}
Model B_sentiment_hybrid: target=resid_PP, regressors=['ret_CRUDE', 'ret_PGP',
'index_hybrid'], target_lags=[1], reg_lags={'ret_CRUDE': [1], 'ret_PGP': [1],
'index_hybrid': [0, 1]}
Model C_sentiment_sections: target=resid_PP, regressors=['ret_CRUDE', 'ret_PGP',
'index_hybrid', 'index_demand', 'index_supply', 'index_price_outlook'],
target_lags=[1], reg_lags={'ret_CRUDE': [1], 'ret_PGP': [1], 'index_hybrid': [0,
1], 'index_demand': [0, 1], 'index_supply': [0, 1], 'index_price_outlook': [0,
1]}
```

```
[4]: backtest_df = pd.DataFrame()

if df_model.empty:
    print("Skipping backtest; dataset is empty.")
elif not specs:
    print("Skipping backtest; no model specs.")
else:
    TEST_MONTHS = 24
    MIN_TRAIN = 36
    backtest_df = run_models_backtest(df_model, specs,
→test_size_months=TEST_MONTHS, min_train_months=MIN_TRAIN)
    if backtest_df.empty:
        print("Backtest produced no results; check data length or specs.")
    else:
        print(f"Backtest rows: {len(backtest_df)}, models: {backtest_df['model_name'].nunique()}")
        display(backtest_df.head())
        out_path = ART / "sentiment_backtest_results.csv"
        backtest_df.to_csv(out_path, index=False)
        print(f"Saved backtest results -> {out_path}")
```

```
Backtest rows: 72, models: 3
```

date	y_true	y_pred	model_name	train_end
------	--------	--------	------------	-----------

```

0 2023-04-30 -0.014802 -0.005182 A_baseline 2023-03-31
1 2023-05-31 -0.053684 -0.002118 A_baseline 2023-04-30
2 2023-06-30  0.000785 -0.013949 A_baseline 2023-05-31
3 2023-07-31  0.016094 -0.000106 A_baseline 2023-06-30
4 2023-08-31  0.049295  0.005123 A_baseline 2023-07-31

```

Saved backtest results -> c:\PythonProjects\LLM-polypropylene\artifacts\sentiment\_backtest\_results.csv

```
[5]: metrics_df = pd.DataFrame()
if backtest_df.empty:
    print("Skipping metrics; no backtest results.")
else:
    metrics_df = compute_forecast_metrics(backtest_df)
    display(metrics_df)
    metrics_path = ART / "sentiment_model_metrics.csv"
    metrics_df.to_csv(metrics_path, index=False)
    print(f"Saved metrics -> {metrics_path}")
```

	model_name	N	MAE	RMSE	MAPE	\
0	A_baseline	24	0.015783	0.020950	0.735172	
1	B_sentiment_hybrid	24	0.018708	0.022892	0.558053	
2	C_sentiment_sections	24	0.029359	0.034768	-0.083298	

	Directional_Accuracy
0	0.458333
1	0.375000
2	0.375000

Saved metrics -> c:\PythonProjects\LLM-polypropylene\artifacts\sentiment\_model\_metrics.csv

```
[6]: coefs_df = pd.DataFrame()
if df_model.empty or not specs:
    print("Skipping coefficient extraction; missing data/specs.")
else:
    coefs_df = extract_coefficients(df_model, specs)
    display(coefs_df.head())
    coefs_path = ART / "sentiment_model_coefficients.csv"
    coefs_df.to_csv(coefs_path, index=False)
    print(f"Saved coefficients -> {coefs_path}")
```

	model_name	variable	coefficient
0	A_baseline	resid_PP_lag1	-0.227014
1	A_baseline	ret_CRUDE_lag1	-0.029404
2	A_baseline	ret_PGP_lag1	0.454791
3	A_baseline	intercept	0.001219
4	B_sentiment_hybrid	resid_PP_lag1	0.269924

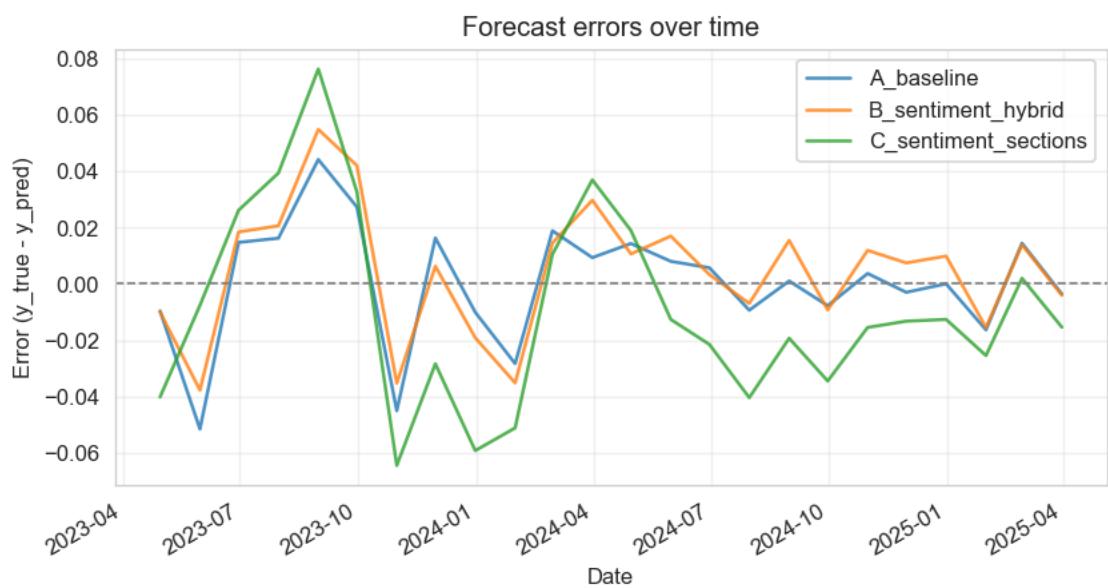
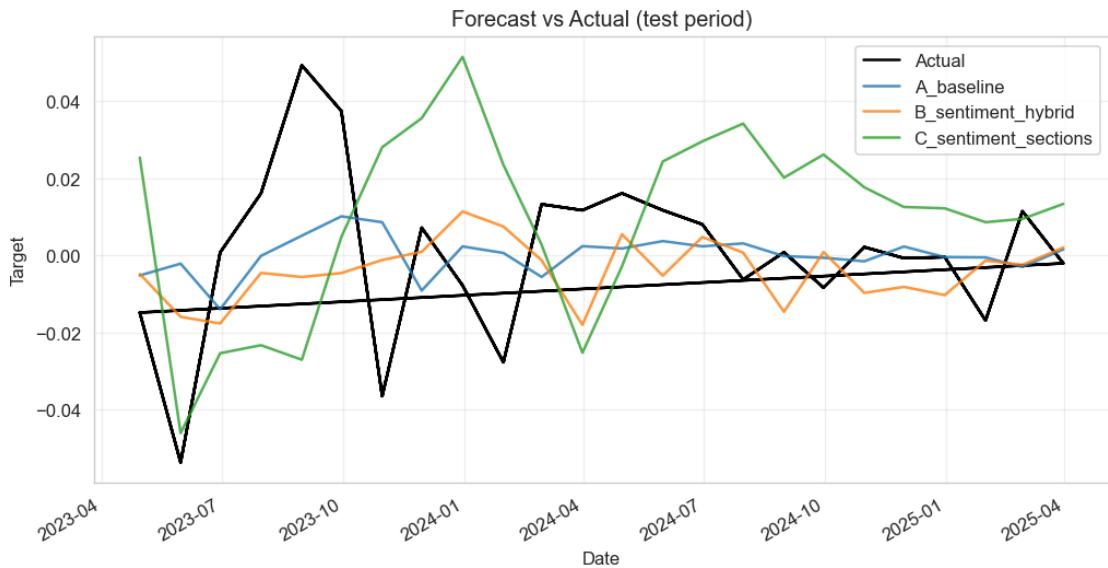
Saved coefficients -> c:\PythonProjects\LLM-

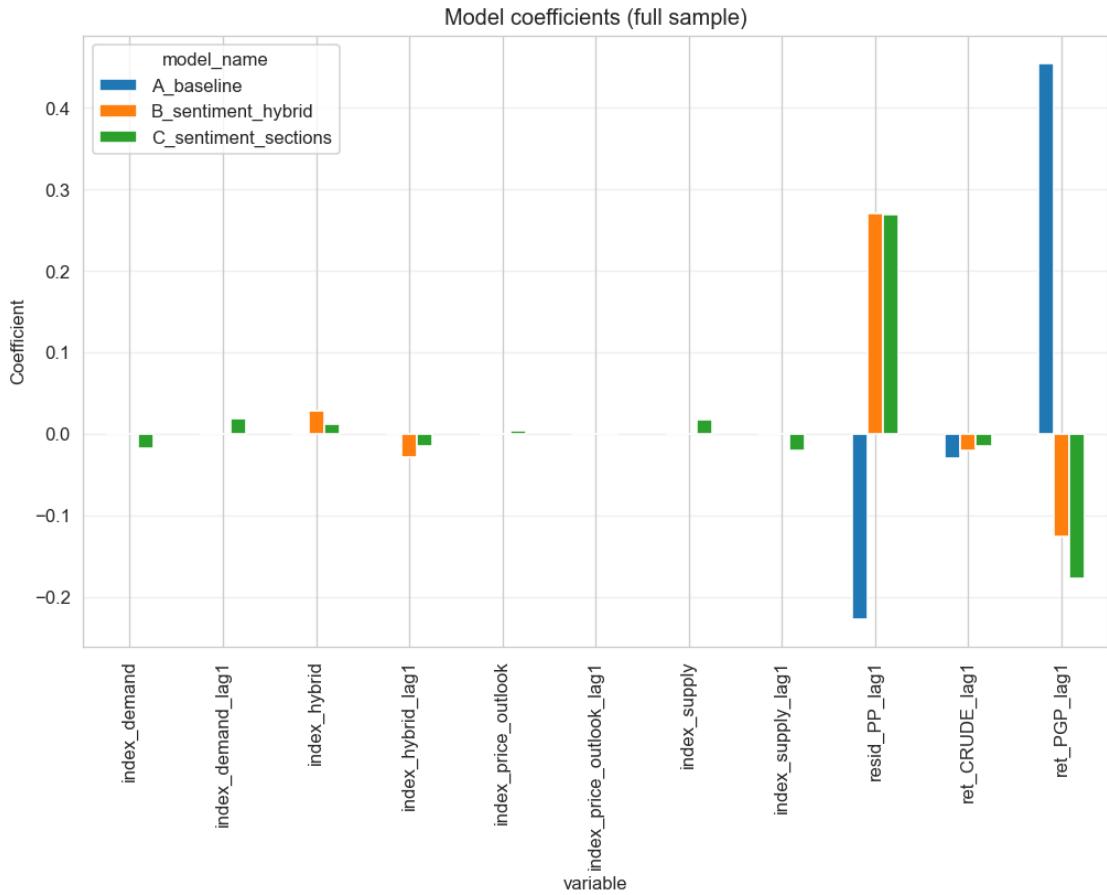
polypropylene\artifacts\sentiment\_model\_coefficients.csv

```
[7]: if backtest_df.empty:
    print("Skipping plots; no backtest results.")
else:
    try:
        fig1 = plot_forecast_vs_actual(backtest_df)
        path1 = PLOTS / "forecast_vs_actual_baseline_vs_sentiment.png"
        fig1.savefig(path1, dpi=150, bbox_inches="tight")
        display(fig1)
    except Exception as exc:
        print(f"Forecast vs actual plot failed: {exc}")
    finally:
        plt.close('all')

    try:
        fig2 = plot_forecast_errors(backtest_df)
        path2 = PLOTS / "forecast_errors_baseline_vs_sentiment.png"
        fig2.savefig(path2, dpi=150, bbox_inches="tight")
        display(fig2)
    except Exception as exc:
        print(f"Forecast errors plot failed: {exc}")
    finally:
        plt.close('all')

if not coefs_df.empty:
    try:
        fig3 = plot_coefficients(coefs_df)
        path3 = PLOTS / "sentiment_model_coefficients.png"
        fig3.savefig(path3, dpi=150, bbox_inches="tight")
        display(fig3)
    except Exception as exc:
        print(f"Coefficient plot failed: {exc}")
    finally:
        plt.close('all')
```





```
[8]: if metrics_df is None or metrics_df.empty:
    print("No metrics to interpret.")
else:
    best = metrics_df.sort_values('RMSE').iloc[0]
    print(f"Best model by RMSE: {best['model_name']} (RMSE={best['RMSE']:.4f}, MAE={best['MAE']:.4f})")
    if 'A_baseline' in metrics_df['model_name'].values:
        base_rmse = metrics_df.loc[metrics_df['model_name']=='A_baseline', 'RMSE'].iloc[0]
        if base_rmse > 0:
            improvement = (base_rmse - best['RMSE']) / base_rmse
            print(f"Improvement vs baseline: {improvement*100:.1f}%")
    print("Use these results to discuss whether sentiment adds forecasting value.")
```

Best model by RMSE: A\_baseline (RMSE=0.0209, MAE=0.0158)

Improvement vs baseline: 0.0%

Use these results to discuss whether sentiment adds forecasting value.